

```

import numpy as np
import pandas as pd

df = pd.read_csv('spam.csv',encoding= "latin-1")

df.sample(5)

```

	v1	v2	Unnamed: 3	Unnamed: 4
2 \				
612	ham	I have many dependents	NaN	NaN
4350	ham	Night has ended for another day, morning has c...	NaN	NaN
4456	ham	Aight should I just plan to come up later toni...	NaN	NaN
4486	ham	Miss call miss call khelate kintu opponenter m...	NaN	NaN
789	ham	Gud mrng dear hav a nice day	NaN	NaN

```

df.shape
(5572, 5)

```

1. Data Cleaning

```

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   v1              5572 non-null   object
 1   v2              5572 non-null   object
 2   Unnamed: 2      50 non-null     object
 3   Unnamed: 3      12 non-null     object
 4   Unnamed: 4      6 non-null      object
dtypes: object(5)
memory usage: 217.8+ KB

df.sample(5)

```

```

v1
2 \
5466 spam http//tms. widelive.com/index. wml?id=820554ad...
NaN
1212 ham Yo, the game almost over? Want to go to walmar...
NaN
2958 ham Buzzzz! *grins* Did I buzz your ass? Buzz your...
NaN
7 ham As per your request 'Melle Melle (Oru Minnamin...
NaN
4591 ham Right it wasnt you who phoned it was someone w...
NaN

```

```

Unnamed: 3 Unnamed: 4
5466 NaN NaN
1212 NaN NaN
2958 NaN NaN
7 NaN NaN
4591 NaN NaN

```

renaming the cols

```

df.rename(columns={'v1':'target','v2':'text'},inplace=True)
df.sample(5)

```

```

target text
Unnamed: 2 \
1394 ham R we still meeting 4 dinner tonight?
NaN
1906 ham And stop being an old man. You get to build sn...
NaN
4016 ham Eek that's a lot of time especially since Amer...
NaN
922 ham It shall be fine. I have avalarr now. Will hol...
NaN
5244 ham thanks for the temales it was wonderful. Thank...
NaN

```

```

Unnamed: 3 Unnamed: 4
1394 NaN NaN
1906 NaN NaN
4016 NaN NaN
922 NaN NaN
5244 NaN NaN

```

```

from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()

```

```

df['target'] = encoder.fit_transform(df['target'])
df.head()

```

```

    target                                text Unnamed:
2  \
0      0  Go until jurong point, crazy.. Available only ...
NaN
1      0                                Ok lar... Joking wif u oni...
NaN
2      1  Free entry in 2 a wkly comp to win FA Cup fina...
NaN
3      0  U dun say so early hor... U c already then say...
NaN
4      0  Nah I don't think he goes to usf, he lives aro...
NaN

```

```

    Unnamed: 3 Unnamed: 4
0      NaN      NaN
1      NaN      NaN
2      NaN      NaN
3      NaN      NaN
4      NaN      NaN

```

missing values

```
df.isnull().sum()
```

```

target      0
text        0
Unnamed: 2   5522
Unnamed: 3   5560
Unnamed: 4   5566
dtype: int64

```

check for duplicate values

```
df.duplicated().sum()
```

```
403
```

remove duplicates

```
df = df.drop_duplicates(keep='first')
```

```
df.duplicated().sum()
```

```
0
```

```
df.shape
```

```
(5169, 5)
```

2.EDA

```
df.head()
```

```

    target                                text Unnamed:
2  \

```

```

0      0  Go until jurong point, crazy.. Available only ...
NaN
1      0                               Ok lar... Joking wif u oni...
NaN
2      1  Free entry in 2 a wkly comp to win FA Cup fina...
NaN
3      0  U dun say so early hor... U c already then say...
NaN
4      0  Nah I don't think he goes to usf, he lives aro...
NaN

```

```

      Unnamed: 3 Unnamed: 4
0      NaN      NaN
1      NaN      NaN
2      NaN      NaN
3      NaN      NaN
4      NaN      NaN

```

```
df['target'].value_counts()
```

```


0      4516
1       653
Name: target, dtype: int64

```

```

import matplotlib.pyplot as plt
plt.pie(df['target'].value_counts(),
labels=['ham', 'spam'], autopct="%0.2f")
plt.show()

```



Data is imbalanced

```
import nltk
```

```
!pip install nltk
```

```
Requirement already satisfied: nltk in c:\users\dell\anaconda3\lib\
site-packages (3.6.5)
Requirement already satisfied: click in c:\users\dell\anaconda3\lib\
site-packages (from nltk) (8.0.3)
Requirement already satisfied: joblib in c:\users\dell\anaconda3\lib\
site-packages (from nltk) (1.1.0)
Requirement already satisfied: regex<=2021.8.3 in c:\users\dell\
anaconda3\lib\site-packages (from nltk) (2021.8.3)
Requirement already satisfied: tqdm in c:\users\dell\anaconda3\lib\
site-packages (from nltk) (4.62.3)
Requirement already satisfied: colorama in c:\users\dell\anaconda3\
lib\site-packages (from click->nltk) (0.4.4)
```

```
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\Dell\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

```
True
```

```
df['num_characters'] = df['text'].apply(len)
```

```
df.head()
```

```
target text Unnamed:
2 \
0 0 Go until jurong point, crazy.. Available only ...
NaN
1 0 Ok lar... Joking wif u oni...
NaN
2 1 Free entry in 2 a wkly comp to win FA Cup fina...
NaN
3 0 U dun say so early hor... U c already then say...
NaN
4 0 Nah I don't think he goes to usf, he lives aro...
NaN
```

```
Unnamed: 3 Unnamed: 4 num_characters
0 NaN NaN 111
1 NaN NaN 29
2 NaN NaN 155
3 NaN NaN 49
4 NaN NaN 61
```

```
# num of words
```

```
df['num_words'] = df['text'].apply(lambda
x:len(nltk.word_tokenize(x)))
```

```
df.head()
```

```

    target                                     text Unnamed:
2  \
0      0  Go until jurong point, crazy.. Available only ...
NaN
1      0                                     Ok lar... Joking wif u oni...
NaN
2      1  Free entry in 2 a wkly comp to win FA Cup fina...
NaN
3      0  U dun say so early hor... U c already then say...
NaN
4      0  Nah I don't think he goes to usf, he lives aro...
NaN

```

```

    Unnamed: 3 Unnamed: 4 num_characters num_words
0      NaN      NaN          111          24
1      NaN      NaN           29           8
2      NaN      NaN          155          37
3      NaN      NaN           49          13
4      NaN      NaN           61          15

```

```

df['num_sentences'] = df['text'].apply(lambda
x:len(nltk.sent_tokenize(x)))

```

```

df.head()

```

```

    target                                     text Unnamed:
2  \
0      0  Go until jurong point, crazy.. Available only ...
NaN
1      0                                     Ok lar... Joking wif u oni...
NaN
2      1  Free entry in 2 a wkly comp to win FA Cup fina...
NaN
3      0  U dun say so early hor... U c already then say...
NaN
4      0  Nah I don't think he goes to usf, he lives aro...
NaN

```

```

    Unnamed: 3 Unnamed: 4 num_characters num_words num_sentences
0      NaN      NaN          111          24              2
1      NaN      NaN           29           8              2
2      NaN      NaN          155          37              2
3      NaN      NaN           49          13              1
4      NaN      NaN           61          15              1

```

```

df[['num_characters', 'num_words', 'num_sentences']].describe()

```

```

    num_characters num_words num_sentences
count    5169.000000  5169.000000  5169.000000
mean       78.977945   18.455407    1.961308
std       58.236293   13.322448    1.432583

```

min	2.000000	1.000000	1.000000
25%	36.000000	9.000000	1.000000
50%	60.000000	15.000000	1.000000
75%	117.000000	26.000000	2.000000
max	910.000000	220.000000	38.000000

ham

```
df[df['target'] == 0]
[['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences
count	4516.000000	4516.000000	4516.000000
mean	70.459256	17.123339	1.815545
std	56.358207	13.491315	1.364098
min	2.000000	1.000000	1.000000
25%	34.000000	8.000000	1.000000
50%	52.000000	13.000000	1.000000
75%	90.000000	22.000000	2.000000
max	910.000000	220.000000	38.000000

#spam


```
df[df['target'] == 1]
[['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences
count	653.000000	653.000000	653.000000
mean	137.891271	27.667688	2.969372
std	30.137753	7.008418	1.488910
min	13.000000	2.000000	1.000000
25%	132.000000	25.000000	2.000000
50%	149.000000	29.000000	3.000000
75%	157.000000	32.000000	4.000000
max	224.000000	46.000000	9.000000

```
import seaborn as sns
```


```
plt.figure(figsize=(12,6))
sns.histplot(df[df['target'] == 0]['num_characters'])
sns.histplot(df[df['target'] == 1]['num_characters'],color='red')
```

```
<AxesSubplot:xlabel='num_characters', ylabel='Count'>
```




```
plt.figure(figsize=(12,6))  
sns.histplot(df[df['target'] == 0]['num_words'])  
sns.histplot(df[df['target'] == 1]['num_words'],color='red')
```

<AxesSubplot:xlabel='num_words', ylabel='Count'>



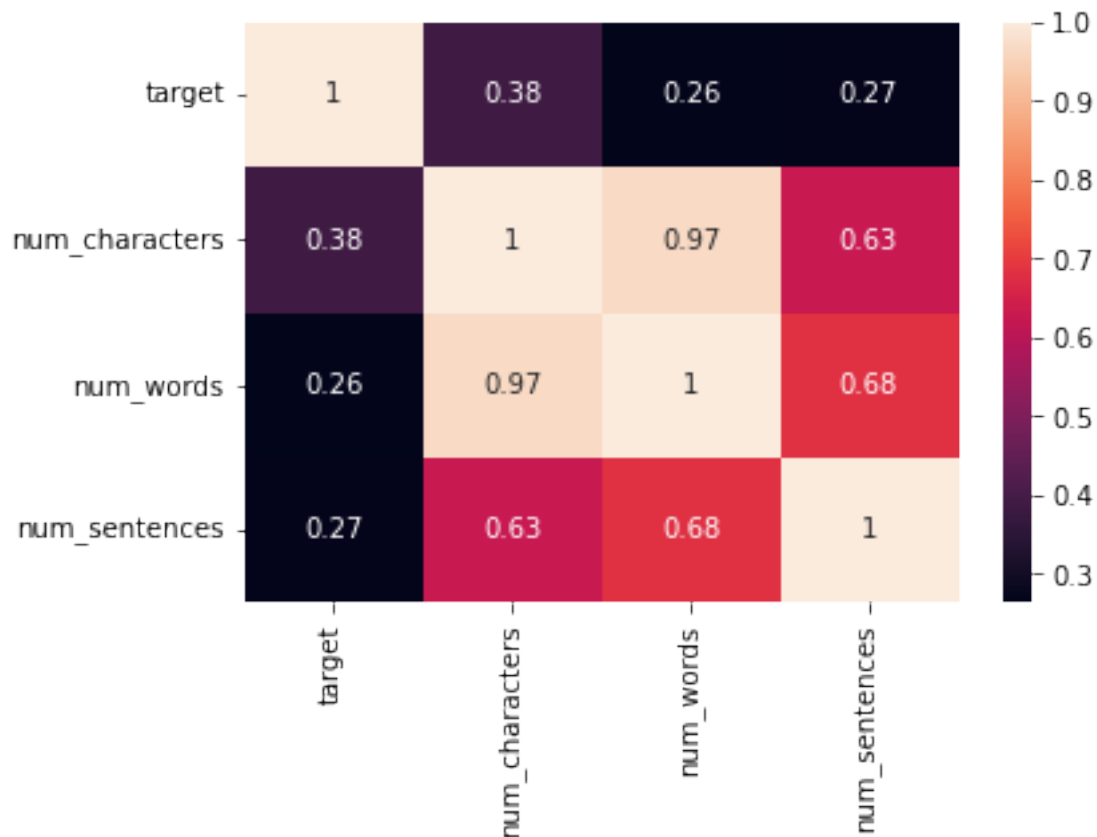
```
sns.pairplot(df,hue='target')
```

<seaborn.axisgrid.PairGrid at 0x1f7d96d0e80>



```
sns.heatmap(df.corr(),annot=True)
```

```
<AxesSubplot:>
```



3. Data Preprocessing

```
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords

import string
string.punctuation

from nltk.stem.porter import PorterStemmer
ps = PorterStemmer()

def transform_text(text):
    text = text.lower()
    text = nltk.word_tokenize(text)

    y=[]
    for i in text:
        if i.isalnum():
            y.append(i)

    text = y[:]
    y.clear()
```

```

    for i in text:
        if i not in stopwords.words('english') and i not in
string.punctuation:
            y.append(i)

```

```

text = y[:]
y.clear

```

```

    for i in text:
        y.append(ps.stem(i))
    return " ".join(y)

```

```

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Dell\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!

```

```

transform_text("I'm gonna be home soon and i don't want to talk about
this stuff anymore tonight, k? I've cried enough today.")

```

```

'gon na home soon want talk stuff anymore tonight k cried enough today
gon na home soon want talk stuff anymor tonight k cri enough today'

```

```

df['text'][10]

```

```

"I'm gonna be home soon and i don't want to talk about this stuff
anymore tonight, k? I've cried enough today."

```

```

from nltk.stem.porter import PorterStemmer
ps = PorterStemmer()
ps.stem('loving')

```

```

'love'

```

```

df['transformed_text'] = df['text'].apply(transform_text)

```

```

df.head()

```

	target	text	Unnamed:
2	\		
0	0	Go until jurong point, crazy.. Available only ...	
NaN			
1	0	Ok lar... Joking wif u oni...	
NaN			
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	
NaN			
3	0	U dun say so early hor... U c already then say...	
NaN			
4	0	Nah I don't think he goes to usf, he lives aro...	
NaN			

	Unnamed: 3	Unnamed: 4	num_characters	num_words	num_sentences	\
0	NaN	NaN	111	24	2	
1	NaN	NaN	29	8	2	
2	NaN	NaN	155	37	2	
3	NaN	NaN	49	13	1	
4	NaN	NaN	61	15	1	

	transformed_text
0	go jurong point crazy available bugis n great ...
1	ok lar joking wif u oni ok lar joke wif u oni
2	free entry 2 wkly comp win fa cup final tkts 2...
3	u dun say early hor u c already say u dun say ...
4	nah think goes usf lives around though nah thi...


```
spam_corpus = []
for msg in df[df['target'] == 1]['transformed_text'].tolist():
    for word in msg.split():
        spam_corpus.append(word)
```

```
len(spam_corpus)
```

```
19878
```

```
from collections import Counter
sns.barplot(pd.DataFrame(Counter(spam_corpus).most_common(30))
[0],pd.DataFrame(Counter(spam_corpus).most_common(30))[1])
plt.xticks(rotation='vertical')
plt.show()
```

```
C:\Users\Dell\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.
warnings.warn(
```




```
ham_corpus = []
for msg in df[df['target'] == 0]['transformed_text'].tolist():
    for word in msg.split():
        ham_corpus.append(word)
```

```
len(ham_corpus)
```

```
70804
```

```
from collections import Counter
sns.barplot(pd.DataFrame(Counter(ham_corpus).most_common(30))
[0],pd.DataFrame(Counter(ham_corpus).most_common(30))[1])
plt.xticks(rotation='vertical')
plt.show()
```

C:\Users\Dell\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.
warnings.warn(



```
df.head()
```

	target	text	Unnamed:
2	\		
0	0	Go until jurong point, crazy.. Available only ...	
NaN			
1	0	Ok lar... Joking wif u oni...	
NaN			
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	
NaN			
3	0	U dun say so early hor... U c already then say...	
NaN			
4	0	Nah I don't think he goes to usf, he lives aro...	
NaN			

	Unnamed: 3	Unnamed: 4	num_characters	num_words	num_sentences	\
0	NaN	NaN	111	24	2	
1	NaN	NaN	29	8	2	
2	NaN	NaN	155	37	2	
3	NaN	NaN	49	13	1	
4	NaN	NaN	61	15	1	

	transformed_text
0	go jurong point crazy available bugis n great ...
1	ok lar joking wif u oni ok lar joke wif u oni
2	free entry 2 wkly comp win fa cup final tkts 2...
3	u dun say early hor u c already say u dun say ...
4	nah think goes usf lives around though nah thi...

4. Model Building

```
from sklearn.feature_extraction.text import
CountVectorizer,TfidfVectorizer
cv = CountVectorizer()
tfidf = TfidfVectorizer(max_features=3000)

X = tfidf.fit_transform(df['transformed_text']).toarray()

X.shape

(5169, 3000)

y = df['target'].values

from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.2,random_state=2)

from sklearn.naive_bayes import GaussianNB,MultinomialNB,BernoulliNB
from sklearn.metrics import
accuracy_score,confusion_matrix,precision_score

gnb = GaussianNB()
mnb = MultinomialNB()
bnb = BernoulliNB()

gnb.fit(X_train,y_train)
y_pred1 = gnb.predict(X_test)
print(accuracy_score(y_test,y_pred1))
print(confusion_matrix(y_test,y_pred1))
print(precision_score(y_test,y_pred1))

0.874274661508704
[[787 109]
 [ 21 117]]
0.5176991150442478

mnb.fit(X_train,y_train)
y_pred2 = mnb.predict(X_test)
print(accuracy_score(y_test,y_pred2))
print(confusion_matrix(y_test,y_pred2))
print(precision_score(y_test,y_pred2))

0.9748549323017408
[[896   0]
 [ 26 112]]
1.0

bnb.fit(X_train,y_train)
y_pred3 = bnb.predict(X_test)
print(accuracy_score(y_test,y_pred3))
```

```
print(confusion_matrix(y_test,y_pred3))
print(precision_score(y_test,y_pred3))
```

```
0.9825918762088974
```

```
[[895   1]
 [ 17 121]]
```

```
0.9918032786885246
```

```
# tfidf --> MNB
```

```
pip install xgboost
```

```
Collecting xgboost
```

```
  Downloading xgboost-1.5.2-py3-none-win_amd64.whl (106.6 MB)
```

```
Requirement already satisfied: numpy in c:\users\dell\anaconda3\lib\
site-packages (from xgboost) (1.20.3)
```

```
Requirement already satisfied: scipy in c:\users\dell\anaconda3\lib\
site-packages (from xgboost) (1.7.1)
```

```
Installing collected packages: xgboost
```

```
Successfully installed xgboost-1.5.2
```

```
Note: you may need to restart the kernel to use updated packages.
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
```

```
svc = SVC(kernel='sigmoid', gamma=1.0)
knc = KNeighborsClassifier()
mnb = MultinomialNB()
dtc = DecisionTreeClassifier(max_depth=5)
lrc = LogisticRegression(solver='liblinear', penalty='l1')
rfc = RandomForestClassifier(n_estimators=50, random_state=2)
abc = AdaBoostClassifier(n_estimators=50, random_state=2)
bc = BaggingClassifier(n_estimators=50, random_state=2)
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
gbdt = GradientBoostingClassifier(n_estimators=50, random_state=2)
xgb = XGBClassifier(n_estimators=50, random_state=2)
```

```
clfs = {
    'SVC' : svc,
    'KN' : knc,
    'NB' : mnb,
    'DT' : dtc,
    'LR' : lrc,
```



```

    'RF': rfc,
    'AdaBoost': abc,
    'BgC': bc,
    'ETC': etc,
    'GBDT': gbd,
    'xgb': xgb
}

def train_classifier(clf, X_train, y_train, X_test, y_test):
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)

    return accuracy, precision

train_classifier(svc, X_train, y_train, X_test, y_test)

(0.9748549323017408, 0.9745762711864406)

accuracy_scores = []
precision_scores = []

for name, clf in clfs.items():

    current_accuracy, current_precision = train_classifier(clf,
X_train, y_train, X_test, y_test)

    print("For ", name)
    print("Accuracy - ", current_accuracy)
    print("Precision - ", current_precision)

    accuracy_scores.append(current_accuracy)
    precision_scores.append(current_precision)

For SVC
Accuracy - 0.9748549323017408
Precision - 0.9745762711864406
For KN
Accuracy - 0.9061895551257253
Precision - 1.0
For NB
Accuracy - 0.9748549323017408
Precision - 1.0
For DT
Accuracy - 0.937137330754352
Precision - 0.8543689320388349
For LR
Accuracy - 0.9545454545454546
Precision - 0.9595959595959596
For RF

```

```

Accuracy - 0.97678916827853
Precision - 0.9830508474576272
For AdaBoost
Accuracy - 0.9680851063829787
Precision - 0.9487179487179487
For BgC
Accuracy - 0.9613152804642167
Precision - 0.8951612903225806
For ETC
Accuracy - 0.9748549323017408
Precision - 0.9590163934426229
For GBDT
Accuracy - 0.9477756286266924
Precision - 0.9468085106382979

```

```

C:\Users\Dell\anaconda3\lib\site-packages\xgboost\sklearn.py:1224:
UserWarning: The use of label encoder in XGBClassifier is deprecated
and will be removed in a future release. To remove this warning, do
the following: 1) Pass option use_label_encoder=False when
constructing XGBClassifier object; and 2) Encode your labels (y) as
integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
  warnings.warn(label_encoder_deprecation_msg, UserWarning)

```

```

[16:40:33] WARNING: C:/Users/Administrator/workspace/xgboost-
win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0,
the default evaluation metric used with the objective
'binary:logistic' was changed from 'error' to 'logloss'. Explicitly
set eval_metric if you'd like to restore the old behavior.
For xgb
Accuracy - 0.9661508704061895
Precision - 0.9557522123893806

```

```

performance_df =
pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy':accuracy_scores,'Prec
ision':precision_scores}).sort_values('Precision',ascending=False)

```

```
performance_df
```


	Algorithm	Accuracy	Precision
1	KN	0.906190	1.000000
2	NB	0.974855	1.000000
5	RF	0.976789	0.983051
0	SVC	0.974855	0.974576
4	LR	0.954545	0.959596
8	ETC	0.974855	0.959016
10	xgb	0.966151	0.955752
6	AdaBoost	0.968085	0.948718
9	GBDT	0.947776	0.946809
7	BgC	0.961315	0.895161
3	DT	0.937137	0.854369

```
performance_df1 = pd.melt(performance_df, id_vars = "Algorithm")
```

```
performance_df1
```

	Algorithm	variable	value
0	KN	Accuracy	0.906190
1	NB	Accuracy	0.974855
2	RF	Accuracy	0.976789
3	SVC	Accuracy	0.974855
4	LR	Accuracy	0.954545
5	ETC	Accuracy	0.974855
6	xgb	Accuracy	0.966151
7	AdaBoost	Accuracy	0.968085
8	GBDT	Accuracy	0.947776
9	BgC	Accuracy	0.961315
10	DT	Accuracy	0.937137
11	KN	Precision	1.000000
12	NB	Precision	1.000000
13	RF	Precision	0.983051
14	SVC	Precision	0.974576
15	LR	Precision	0.959596
16	ETC	Precision	0.959016
17	xgb	Precision	0.955752
18	AdaBoost	Precision	0.948718
19	GBDT	Precision	0.946809
20	BgC	Precision	0.895161
21	DT	Precision	0.854369

```
sns.catplot(x = 'Algorithm', y='value',  
            hue = 'variable',data=performance_df1,  
            kind='bar',height=5)  
plt.ylim(0.5,1.0)  
plt.xticks(rotation='vertical')  
plt.show()
```



```
# model improve
# 1. Change the max_features parameter of TfIdf

temp_df =
pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_max_ft_3000':accuracy_
scores,'Precision_max_ft_3000':precision_scores}).sort_values('Precisi
on_max_ft_3000',ascending=False)

temp_df =
pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_scaling':accuracy_scor
es,'Precision_scaling':precision_scores}).sort_values('Precision_scali
ng',ascending=False)

new_df = performance_df.merge(temp_df,on='Algorithm')

new_df_scaled = new_df.merge(temp_df,on='Algorithm')

temp_df =
pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_num_chars':accuracy_sc
```

```
ores, 'Precision_num_chars':precision_scores}).sort_values('Precision_num_chars',ascending=False)
```

```
new_df_scaled.merge(temp_df,on='Algorithm')
```

	Algorithm	Accuracy	Precision	Accuracy_scaling_x
0	KN	0.906190	1.000000	0.906190
1	NB	0.974855	1.000000	0.974855
2	RF	0.976789	0.983051	0.976789
3	SVC	0.974855	0.974576	0.974855
4	LR	0.954545	0.959596	0.954545
5	ETC	0.974855	0.959016	0.974855
6	xgb	0.966151	0.955752	0.966151
7	AdaBoost	0.968085	0.948718	0.968085
8	GBDT	0.947776	0.946809	0.947776
9	BgC	0.961315	0.895161	0.961315
10	DT	0.937137	0.854369	0.937137

	Accuracy_scaling_y	Precision_scaling_y	Accuracy_num_chars
0	0.906190	1.000000	0.906190
1	0.974855	1.000000	0.974855
2	0.976789	0.983051	0.976789
3	0.974855	0.974576	0.974855
4	0.954545	0.959596	0.954545
5	0.974855	0.959016	0.974855
6	0.966151	0.955752	0.966151
7	0.968085	0.948718	0.968085
8	0.947776	0.946809	0.947776
9	0.961315	0.895161	0.961315
10	0.937137	0.854369	0.937137

	Precision_num_chars
0	1.000000
1	1.000000
2	0.983051
3	0.974576
4	0.959596
5	0.959016

```
6          0.955752
7          0.948718
8          0.946809
9          0.895161
10         0.854369
```

Voting Classifier

```
svc = SVC(kernel='sigmoid', gamma=1.0, probability=True)
mnb = MultinomialNB()
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)

from sklearn.ensemble import VotingClassifier

voting = VotingClassifier(estimators=[('svm', svc), ('nb', mnb),
('et', etc)], voting='soft')

voting.fit(X_train, y_train)

VotingClassifier(estimators=[('svm',
                             SVC(gamma=1.0, kernel='sigmoid',
                                probability=True)),
                             ('nb', MultinomialNB()),
                             ('et',
                              ExtraTreesClassifier(n_estimators=50,
                                                    random_state=2))],
                 voting='soft')

y_pred = voting.predict(X_test)
print("Accuracy", accuracy_score(y_test, y_pred))
print("Precision", precision_score(y_test, y_pred))

Accuracy 0.9796905222437138
Precision 0.9834710743801653

# Applying stacking
estimators=[('svm', svc), ('nb', mnb), ('et', etc)]
final_estimator=RandomForestClassifier()

from sklearn.ensemble import StackingClassifier

clf = StackingClassifier(estimators=estimators,
                        final_estimator=final_estimator)

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print("Accuracy", accuracy_score(y_test, y_pred))
print("Precision", precision_score(y_test, y_pred))

import pickle
pickle.dump(tfidf, open('vectorizer.pkl', 'wb'))
pickle.dump(mnb, open('model.pkl', 'wb'))
```